

AVERY CHONG

+1 672 866 2568 | averycx14@gmail.com | linkedin.com/in/cxinyu | github.com/nwfish

EDUCATION

University of British Columbia

Bachelor of Science, Major in Computer Science

Vancouver, Canada

Expected May 2029

- **Coursework:** Applied Machine Learning, Data Science, Statistics, Algorithms & Data Structures, Relational Databases, Intro to Software Engineering, Relational Databases, Operating Systems
- **Awards:** Outstanding International Student Award; Diversity in CS Winner, YouCode Hackathon

TECHNICAL SKILLS

AI & ML: PyTorch, scikit-learn, Hugging Face Transformers, sentence-transformers, Ollama, Groq, Gemini API

Models: Llama 3.1 (8B), DeBERTa-v3, RoBERTa, all-MiniLM-L6-v2, Nomic Embed v1.5

Data Science: pandas, NumPy, Jupyter, matplotlib, Streamlit

Data Engineering: ETL pipelines, Airflow, Docker, CI/CD (GitHub Actions), PostgreSQL, Supabase, MongoDB, Oracle

Languages & Tools: Python, SQL, TypeScript, Java, C/C++, R; FastAPI, Node.js, Express, Git

PROJECTS

Feedle (Photo Feed Curation) | *PyTorch, CLIP, scikit-learn, OpenCV, FastAPI*

Sept 2026

- Building a pipeline that selects and arranges a folder of photos into a cohesive feed, using **CLIP**-based aesthetic scoring in **PyTorch**, OpenCV color and density features, and a logistic regression trained on hand-labeled feeds.

Checkup (Clinical Note Faithfulness Eval) | *Python, Ollama (Llama 3.1 8B), pandas, Streamlit*

Sept 2026

- Built a **local-LLM** pipeline that checks AI-generated intake notes against the patient interview, labeling each field captured, missed, invented, or wrong; it agreed with hand labels on **45/50 (90%)** before tuning.
- Hand-labeled a **50-item gold set** from **10** sampled interviews in UBC's synthetic psychiatry dataset (SPIT), and linked **174 transcripts** to source profiles with no patient ID using a unique-match rule.
- Wrote a rule-based **chain-of-thought leak detector** and found patient-model leaks were **~6×** more common after an interviewer leak (**58% vs. 10%**), and built **2 Streamlit apps** for reviewing and labeling.

SoCurious (NLP Question Pipeline) | *Python, Hugging Face, sentence-transformers, PostgreSQL*

Aug 2026

- Benchmarked sentence-embedding deduplication against fuzzy matching on **525 real questions**, and rebuilt the pipeline around embeddings after proving **100% recall vs. 66.7%**.
- Chose a **22M-parameter** embedding model over a **360×** larger alternative after confirming no measurable task improvement, optimizing for inference cost and latency.
- Built a **two-stage classification layer** (zero-shot topic tagging, then a domain-specific scoring model), refined across **5 rounds** of manual validation against real data batches.
- Built a multi-source ETL pipeline (**1,484+ posts, 6 subreddits**) with a **7-rule** filter layer, and caught **2** over-aggressive rules with error analysis before production.

blindspot (YouTube Hate-Speech Pipeline) | *Python, Airflow, Docker, Hugging Face, pandas*

Sept 2026 - Present

- Built an **Airflow** pipeline in **Docker** that ingests YouTube comments into a bronze layer, and scored them for **sentiment and emotion** with Hugging Face models.
- Integrating a **7-class RoBERTa hate classifier** and a **BigQuery star schema** to surface what kinds of hate a creator's community engages in.

Matchie (Diversity in CS Winner) | *TypeScript, Nomic Embed, Groq, MongoDB, Next.js*

YouCode 2026

- Built semantic volunteer matching with **Nomic Embed v1.5** and cosine similarity across **4 MongoDB collections**, replacing brittle rule-based filtering, and added **Groq LLM** match explanations.

EXPERIENCE

nwPlus (UX/UI Designer & External Engagement Coordinator)

Sept 2025 – Present

- Designing website/merch and contributed to **day-of operations** for **HackCamp**, **nwHacks**, and **cmd-f**, reaching **1,250+ participants** as part of a **55-member team**.
- Led **Pitching 101** and **Inclusive Language** workshops for **50+ hackers** and **40+ mentors/judges/volunteers**, and will lead a **database workshop**.